

Idiomatic Modern C++ for Linux

Week 4: Advanced Type Concepts

Today's agenda

- type deduction
- type aliases
- const
- constexpr variables and functions
- conversions and type casting

The 'auto' keyword

- All objects need to be of an explicit type, however, some literals have an implicitly determined type depending on their format.
- The ***auto*** keyword can be used in place of a type specifier to deduce the type of object from its initializer.
- auto can also be used to deduce the return type of a function, in the function definition itself. (cursed)
- Can be a handy tool when it increases readability

The 'auto' keyword

```
// Default literal values  
auto d {4.6}; // default floating point type is a double, NOT float  
auto i {46}; // int  
auto b {true}; // bool  
auto c {'^'}; // char  
auto cstr = "hi, I'm a C string!"; // const char[20]
```

auto dos and don'ts

Do use auto:

- when it increases the readability of your code

```
std::unique_ptr<std::vector<int>> my_vec_ptr =  
    std::make_unique<std::vector<int>>(1, 3, 5, 7);  
  
auto my_vec_ptr = std::make_unique<std::vector<int>>(1, 3, 5, 7);
```

Don't use auto:

- when it obfuscates what's going on with your code
- as the return type for a function, generally speaking

```
auto result = ProcessSomeData(arg);
```

```
// you can use 'auto' to deduce a return type,  
// but there's not a lot of good reasons to do this.  
// if I were a developer using an API that had an 'auto' return type,  
// and needed to dig into the code to figure out a bug I was encountering,  
// I'd probably not be too happy having the return type  
// obscured down further in the function body  
auto mul (int a, int b) {  
    return a * b;  
}
```

Literal suffixes

- C++ lets you specify the type of a literal value with *literal suffixes*
- e.g. “12.4f” is a *float literal*, “12463159ul” is an *unsigned long literal*

Literal suffixes

```
// Literal suffixes  
auto f {6.12f}; // 'f' indicates a float literal.  
auto f2 {8.14F}; // all of these work for upper or lowercase  
auto u {192436U}; // unsigned int  
auto l {16720429612l}; // long int  
auto ll {79611}; // long long int  
auto ull {98235454llu}; // long long unsigned int  
  
// std::string literal.  
// only works with the std::string_literals namespace  
auto str {"hi, I'm a std::string"s };
```

Introducing 2 new stl containers

- `std::pair` – ties two elements together
- Access elements by `pair.first` and `pair.second`
- `std::unordered_map` – stl hashmap implementation.
- Closest analogue to python dictionaries
- Key type must be hashable
- Elements are NOT sorted by key, unlike `std::map`
- $O(1)$ insertion, access and deletion

Type aliases

- Type aliasing lets you create a shorthand for an existing type.
- This can be really convenient when you're using the same verbose type in multiple places and want to cut down on how much code you're writing.
- A type alias can also make the intent of your code much clearer

typedef – C style type aliases

- 'typedef' is C's way of declaring type aliases.
- Like many C features, it still exists in C++ for backwards compatibility
- Some type definitions can be harder to read
- The 'using' keyword has largely replaced typedef and should be preferred

```
typedef std::pair<int, int> Point2i;
```

```
// with typedef, FuncType is buried in the typedef and harder to read  
typedef float (*FuncType)(double, std::string);
```

```
using Point2i = std::pair<int, int>;
```

```
// using makes FuncType much clearer  
using FuncType = float (*)(double, std::string);
```

Nested type aliases

```
std::unordered_map<std::string, std::pair<double, double>> house_locations {  
    {"living_room", {2.3, 5.6}},  
    {"bathroom", {4.1, 4.62}},  
    {"bedroom", {8.4, 9.172}},  
};
```

```
using Point2d = std::pair<double, double>;  
using StringToPointMap = std::unordered_map<std::string, Point2d>;
```

```
StringToPointMap house_locations {  
    {"living_room", {2.3, 5.6}},  
    {"bathroom", {4.1, 4.62}},  
    {"bedroom", {8.4, 9.172}},  
};
```

- you can also create composite type aliases using other type aliases.

about const

- Variables can be declared immutable with the 'const' keyword.
- Values can be determined either at compile-time or runtime, depending on if the value is a constant expression or returned from a function (that isn't constexpr).
- Values declared const are stored in the executable.

```
const double PI = 3.1415926;  
//PI = 3.15; // won't compile, PI declared const  
  
const int random_number = rand(); // determined at runtime  
  
const MyCoolStruct s {100, "gecs"};
```

Const pointers and pointers to const

```
// pointer to a const int.  
// The pointer can point elsewhere,  
// but the data it points to can't be modified  
const int* ptr2const = &random_number;  
ptr2const = &val2;  
//*ptr2const = 4;
```

```
// const pointer to an int.  
// the pointer can't point elsewhere,  
// but the data it points to can be modified  
int* const constptr = &val;  
// constptr = &val2;  
*constptr = 5;
```

Const references

- Const references let you refer to some original value by its location in memory, but prevents mutation and unnecessary copying
- Most of the time, we'll end up using const references when accessing values or passing to a function

```
// When iterating through containers,  
// it's best practice to use const references  
// if you don't need to modify the underlying data.  
std::unordered_map<std::string, unsigned int> some_hashmap {  
    {"foo", 81},  
    {"abc", 123},  
    {"I love strings", 9999999}  
};  
  
// range based for loop, using an unordered map iterator  
for (const auto& it : some_hashmap) {  
    cout << it.first << ", " << it.second << endl;  
}
```

```
void my_cool_function(const MyCoolStruct& my_struct) {  
    my_struct.n = 45; // won't compile, my_struct is a const reference  
}
```

Bonus – Structured bindings

- C++17 introduced **structured bindings**
- Syntactic sugar to make value unpacking easier, and the “cool” for-loops even cooler

```
// structured binding syntax - new in C++17  
for (const auto& [s, n] : some_hashmap) {  
    cout << s << ", " << n << endl;  
}
```

```
const MyCoolStruct s {100, "gecs"};  
auto [hundred, geecs] = s;
```

Introducing constexpr

- The **constexpr** keyword declares that an expression can be evaluated at compile-time.
- This makes it possible to compute expressions and store them *before your program even runs*
- Save CPU cycles by offloading computation to the compiler
- Constexpr can apply to
 - Variables
 - Functions
 - Templates (template metaprogramming)
 - Conditional compilation

Constexpr variable rules

You can initialize a variable as constexpr iff:

- The declaration is also a definition
- It's a literal type
- It's initialized by the declaration
- The full expression of its initialization is constexpr

constexpr variable examples

```
constexpr double TWO_PI = 6.28318530718;
```

```
// Struct aggregates  
// and classes with only constexpr constructors and member methods  
// can also be evaluated at compile-time.
```

```
constexpr MyCustomPoint2iStruct p1 {12, 4};
```

Constexpr function rules

- Non-void return type
- Define with constexpr before the function definition
- No non-const-qualified types in variable definitions, unless they're initialized with constexpr
- Only calls other constexpr functions
- Must produce a constexpr when called with a constexpr

Constexpr function examples

- In order to run at compile time, all args must be known at compile time

```
// fib can evaluate either at compile time or runtime.  
constexpr long fib(long n) {  
    return (n < 2) ? 1 : fib(n-1) + fib(n-2);  
}
```

```
using namespace std::chrono;  
  
auto constexpr_start = high_resolution_clock::now();  
// 24 is a constant literal, so fib will evaluate during compile-time here.  
constexpr long compiletime_fib_result = fib(24);  
auto constexpr_stop = high_resolution_clock::now();  
auto constexpr_duration = duration_cast<nanoseconds> (constexpr_stop - constexpr_start);  
  
auto nonconstexpr_start = high_resolution_clock::now();  
// a was not declared constexpr, so fib will execute at runtime here  
long a {24};  
const long runtime_fib_result = fib(a);  
auto nonconstexpr_stop = high_resolution_clock::now();  
auto nonconstexpr_duration = duration_cast<nanoseconds> (nonconstexpr_stop - nonconstexpr_start);
```

```
nanoseconds took to calculate fib(24) (75025) during compile-time: 29  
nanoseconds took to calculate fib(24) (75025) during runtime: 217467
```

Constexpr function examples

- Functions that can evaluate at compile-time are also great for initializing objects whose size *must be known* at compile time.

```
constexpr int mul (int a, int b) {  
    return a * b;  
}
```

```
constexpr std::array<int, mul(3, 9)> my_arr {1, 2, 3, 4, 5};
```

A word on consteval (C++20 onwards)

- `constexpr` tells the compiler that an expression *may* evaluate at compile-time, but doesn't enforce it.
- ***constexpr*** enforces compile-time evaluation for a function, if for whatever reason you require a function to only ever evaluate at compile-time and not runtime

Conversions – implicit conversion

- Implicit conversion is done automatically by the compiler in contexts where the **expected** type differs from the **actual** type
- Doesn't change the underlying data
- Implicitly involves a temporary copy

Conversions – implicit conversion

- Data is represented by different objects in different ways
- The bits from `n` are copied into `f`, but the underlying bit representation of float differs from int
- the code is certainly not gonna display "51"
- Be careful with `memcpy` as it can result in undefined behavior

```
int n{51};  
float f;  
std::memcpy(&f, &n, sizeof(float));  
cout << "float memcpy'd over to an int: " << f << endl;
```

```
float memcpy'd over to an int: 7.14662e-44
```


Implicit conversion examples

```
// implicitly converts integer literal 3 to a double
double d{3};
// implicitly converts integer literal 6 to a double
d = 6;
```

```
std::function<float()>some_fn([] {
    // double literal implicitly converted to float
    return 3.00000001;
});

float f = some_fn();
```

```
// integer values converted to a boolean
if (17) {
    cout << "integers that are nonzero are implicitly \"true\"" << endl;
}
if (0) {
    cout << "this code shouldn't be reached" << endl;
}
```

```
std::function<void(const long)>some_other_fn([](const long n) {});

int a = 4;
some_other_fn(a); // int implicitly converted to long
```

Implicit conversion can fail

- A conversion will fail if there's no known way to convert a type to another

```
// won't compile. can't convert const char to int  
int bad_convert{"26"};
```

Explicit conversion: C-style casting

- The old C syntax is still present in C++ for backwards compatibility
- There's also "function style" casting syntax present in C++
- C style casting should be avoided in favor of `static_cast`
 - It may perform either a `static_cast`, `const_cast` or `reinterpret_cast`, possibly introducing undefined behavior
 - Harder to read and search for in code
- Only one use case not covered by C++ casts: can convert a derived object to an inaccessible base class (e.g. private inheritance)

```
// old way, C Style cast  
double a = (double)42;  
// alternate "function style" cast  
double b = double(42);
```

Explicit conversion: static_cast

- Explicitly convert a value of one type to a value of another type
- Does compile-time check for if the conversion can be made
- Intentionally less powerful than C style casting, so can't reinterpret a type or cast away constness

```
// Modern way, static_cast (preferred way in C++)  
double c = static_cast<double>(42);
```

Explicit conversion: reinterpret_cast

- Used to reinterpret the underlying *bit pattern* of some type to another type
- Primarily used for pointer conversion
- Can be extremely dangerous and introduce undefined behavior if used incorrectly
- Don't use unless you can't use static_cast or const_cast

```
// Reinterpret cast is primarily used to convert pointer types.  
int n {51};  
  
// We create a pointer to a float, and store the address of d (an integer),  
// but converted as if it were a pointer to a float.  
// You'll notice this yields the same result as our memcpy example from earlier.  
// Reinterpret cast can introduce undefined behavior and should be avoided whenever possible.  
float* f = reinterpret_cast<float*>(&n);  
cout << "pointer to int reinterpret_casted as a pointer to float: " << *f << endl;
```

```
pointer to int reinterpret_casted as a pointer to float: 7.14662e-44
```

reinterpret_cast used in production

- Now that I've just told you to never use `reinterpret_cast`, here's an example of fast square root inverse not unlike the one used in Quake
- Does some wacky pointer reinterpretation combined with one iteration of least squares error minimization

```
float fast_inv_sqrt(float x, int num_gauss_newton_iterations = 1)
{
    float y = x; // current guess y = sqrt(x)

    // i points to current guess y, interpreted as uint32
    uint32_t* i = reinterpret_cast<uint32_t*>(&y);
    constexpr uint32_t exp_mask = 0x7F800000; // 0xFF<<23
    constexpr uint32_t magic_number = 0x5f000000; // 190<<23

    // initial guess using magic number
    *i = magic_number - ((*i >> 1) & exp_mask);

    // refine guess using some Newton iterations
    for (size_t i = 0; i < num_gauss_newton_iterations; ++i)
    {
        y = (x * y * y + 1) / (2 * x * y);
    }

    return y;
}
```

Numeric promotion

- C++ is designed to be portable, and often it's more efficient to process a wider type
- We can just write functions that take an int and double, and “promote” more narrow types, without explicitly defining a function with ‘short’, ‘char’, etc

```
short s{12};  
print_integer(s); // numeric promotion of short to int  
print_integer('a'); // numeric promotion of char to int  
  
// "true" will convert to 1 when promoted to an int  
// whereas "false" converts to 0  
print_integer(true);  
  
print_double(8.6f); // float literal promoted to double  
print_double(96.523124321);
```

Narrowing conversions

- It's best to explicitly `static_cast` to the narrower type if you need to
- Be careful about loss of precision or loss of information

```
double d{3.1415};  
// should be a compiler warning with stricter linting  
// implicit narrowing conversion  
print_integer(d);  
  
// explicitly narrowing, so no compiler warning  
print_integer(static_cast<int>(d));
```


Constexpr and narrowing conversions

- Some constexpr conversions aren't considered narrowing by the compiler.
- Since constexpr asserts that a value is known at compile time, we won't get compiler errors so long as the source value is preserved.

```
int n1{ 5 };  
unsigned int u1 { n1 };
```

```
Build files have been written to: /code-examples/W4-advanced-type-concepts/build  
[ 12%] Building CXX object CMakeFiles/conversion.dir/src/conversion.cpp.o  
/code-examples/W4-advanced-type-concepts/src/conversion.cpp: In function 'void narrowing_conversions()':  
/code-examples/W4-advanced-type-concepts/src/conversion.cpp:131:23: warning: narrowing  
conversion of 'n1' from 'int' to 'unsigned int' [-Wnarrowing]  
 131 |     unsigned int u1 { n1 };  
      |                      ^~  
[ 25%] Linking CXX executable conversion  
[ 25%] Built target conversion
```

```
constexpr int n1{ 5 };  
unsigned int u1 { n1 };
```

```
-- Generating done (0.0s)  
-- Build files have been written to: /code-examples/W4-advanced-type-concepts/build  
[ 25%] Built target conversion
```

Constexpr and narrowing conversions

- Strangely, narrowing a constexpr floating point type is not technically considered narrowing by the compiler, even when there's loss of precision

```
double my_extremely_precise_number {0.426421124085756};  
float totally_not_narrowing{my_extremely_precise_number};  
cout << totally_not_narrowing << endl;
```

```
12%] Building CXX object CMakeFiles/conversion.dir/src/conversion.cpp.o  
code-examples/W4-advanced-type-concepts/src/conversion.cpp: In function 'void narrowin  
_conversions()':  
code-examples/W4-advanced-type-concepts/src/conversion.cpp:141:33: warning: narrowing  
conversion of 'my_extremely_precise_number' from 'double' to 'float' [-Wnarrowing]  
141 |     float totally_not_narrowing{my_extremely_precise_number};  
    |                               ^~  
25%] Linking CXX executable conversion  
25%] Built target conversion
```

```
constexpr double my_extremely_precise_number {0.426421124085756};  
float totally_not_narrowing{my_extremely_precise_number};  
cout << totally_not_narrowing << endl;
```

```
-- Build files have been written to: /code-examples/W4-advanced-type-concepts/build  
[ 12%] Building CXX object CMakeFiles/conversion.dir/src/conversion.cpp.o  
[ 25%] Linking CXX executable conversion
```

Arithmetic conversion

- Certain operators require both operands to have the same type.
- if we invoke one of them with operands of different types, one of the operands will be implicitly converted to match using a set of rules called the usual arithmetic conversions.

Arithmetic conversion

```
// The above code should print out 'i, d' (for integer and double)  
// but what about the following expression?  
auto sum = my_int + my_double;  
cout << typeid(sum).name() << ' ' << sum << endl;
```

```
// Short is not on the priority list for operator+,  
// so both operands get promoted to int.  
short l {9};  
short r {10};  
  
cout << typeid(l + r).name() << endl;
```

Arithmetic conversion

- Things start to get weird when we throw signedness into the mix.
- -3 is promoted to an unsigned integer that ends up being larger than 5

```
cout << std::boolalpha << (-3 < 5u) << endl;
```

false

Additional Resources

- <https://medium.com/@sofiasondh/c-const-vs-constexpr-the-comparison-183f9dd92deb>
- <https://rmminusr.com/cursed-reinterpret-cast/>
- https://en.cppreference.com/w/cpp/language/usual_arithmetic_conversions.html